

Ogham Manual

Full user manual for Ogham — a dual-voice bytebeat synthesizer for Eurorack.

UPDATED 30 JULY 2026

A dual-voice bytebeat synthesizer for Eurorack, built on the Electro-Smith Daisy Seed.

1. Introduction

Ogham is a Eurorack voice that turns tiny mathematical formulas into music. Instead of oscillators and wavetables, it runs bytebeat expressions — single lines of integer arithmetic that, evaluated over time, produce surprisingly rich rhythmic and melodic patterns. Ogham hosts a bank of these formulas, lets you play two of them at once, shape them with a bipolar Tone macro and a 3-stage modulation FX chain, and clock, pitch, and modulate them from the rest of your system.

At a glance

- 100 function slots (**F0** – **F99**), each with two performable parameters (A & B). 22 are authored formulas today; the rest are placeholders (see Section 4), and a special **AA** slot holds a built-in A=440 Hz tuning reference.
 - Two independent voices (Out 1 / Out 2), each playing its own function.
 - Out 2 “drone/freeze” — decouple Out 2 into an independent, self-sustaining voice that ignores the knobs, clock, pitch, and FX.
 - Bipolar Tone macro — one knob from clean through low-pass filtering, wavefolding, saturation, sample-rate reduction, and resonant band-pass grit.
 - 3-stage modulation FX — Chorus → Flanger → Phaser, each with a clean and a characterful variant, in series or parallel.
 - Mode switch: Clk ↔ VOct on a shared jack.
 - Clk — external tempo/clock sets the bytebeat playback rate.
 - VOct — 1 V/oct pitch: the engine hard-syncs to play clean tuned notes.
 - Hard-sync Sync input — sample-accurate waveform reset, usable up to audio rates.
 - CV over A and B (± 5 V, summed with the knobs); either CV can be re-routed to the Tone macro instead of its parameter.
 - Internal LPG — a low-pass gate on the output, plucked from the Sync input: sharp attack, decay from 2 ms to 20 s, filter and VCA moving together.
 - Selectable ENV Out — an envelope follower of either voice, *or* a raw bytebeat used as a slow DC modulation/LFO source; plus an EOC pulse out.
 - Voice selections, the full effects/settings menu, and all options persist across power cycles; a one-gesture factory reset.
-

2. The panel

Ogham packs a lot into a small panel: five knobs (one is the Func encoder), a mode switch, a 4-digit display, four inputs, and four outputs. The map below names every control; the detailed behaviour of each is in Section 5, *Controls — full reference* (the numbers here match the references there).



Control legend

#	Control	Type	In short
1	Func	Rotary encoder (with push)	Select each voice's function; long-press opens the Menu; short-click switches voice
2	A	Knob (+ CV)	Formula Parameter A (0–255), shared by both voices
3	B	Knob (+ CV)	Formula Parameter B (0–255), shared by both voices
4	Rate / Fine	Knob	Playback rate $0.25\times$ – $4\times$ (a quantized clock $\times/\div$ when a clock is patched); ± 1 -octave fine tune in VOct mode

#	Control	Type	In short
5	Tone	Knob	Bipolar tone macro — clean at 12 o'clock, filtering/fold CCW, crush/resonance CW
6	Mode	Toggle switch	Selects the shared jack (11): Clk (tempo) or VOct (pitch)
7	Display	4-digit 7-segment	Current function, menu fields, parameter flashes, clean-center dot
8	CV A	Input (± 5 V)	Modulates Parameter A (summed with knob 2); can be routed to Tone
9	CV B	Input (± 5 V)	Modulates Parameter B (summed with knob 3); can be routed to Tone
10	Sync	Input (trigger/gate)	Hard-sync reset — restarts the waveform (audio-rate capable)
11	Clk/VOct	Input (shared)	Tempo clock or 1 V/oct pitch, per the Mode switch (6)
12	Out 1	Audio output	Voice 1 (primary) — drives the EOC / default envelope
13	Out 2	Audio output	Voice 2 — independent function; can be frozen as a drone
14	ENV Out	CV output (DC-coupled)	Envelope follower of a voice, or a raw bytebeat DC/LFO — mode set in the Menu
15	EOC	Gate output	End-of-cycle pulse (~10 ms), BPM-synced from Voice 1

Reading the display (7): functions show as **1-NN** / **2-NN** (voice + function, 0-based), **1-AA** for the tuning reference, and various codes in the Menu (Section 5). The far-right decimal point lights whenever the Tone knob (5) is at its clean center.

3. What is bytebeat?

Bytebeat is a style of music discovered and popularised by Viznut (Ville-Matias Heikkilä) in 2011. The whole instrument is a single expression:

```
output = f(t)           // f returns an 8-bit value 0-255
```

`t` is a counter that increments at a fixed sample rate (Ogham uses 8 kHz for its authored formulas). You feed `t` through arithmetic and bitwise operators — shifts, AND/OR/XOR, multiply, divide, modulo — and the low 8 bits of the result are sent to the DAC as audio. No envelopes, no note lists; the *structure of the number sequence itself* becomes rhythm, melody, and timbre.

A classic example is `t*(((t>>12)|(t>>8))&(63&(t>>4)))` — a handful of characters that produce an evolving arpeggio (it is, in fact, the placeholder you'll hear in the not-yet-authored slots; see Section 4). Small changes to the constants produce wildly different results, which is exactly what makes bytebeat both fascinating and hard to control.

How Ogham plays it

- Each formula runs at its base rate (8 kHz for the authored bank), then is linearly interpolated up to the 48 kHz codec output for a smoother, less aliased sound.
 - A 32.32 fixed-point phase accumulator bridges the two rates and provides the variable playback rate (the Rate / Fine knob, or an external clock).
 - All maths use JavaScript `ToInt32` semantics (shifts masked to `& 31`, result masked to `& 0xFF`) so formulas behave exactly as they do in the browser bytebeat tools they were designed in.
-

4. The functions — 100 slots

Ogham addresses 100 function slots, shown on the display as **F0** – **F99** (the numbering is 0-based: the first function reads **1-00**). Turning the Func encoder (1) moves through them; the selector does not wrap (see Section 5).

What's in the slots today

Slots	Contents
F0 – F21	The 22 authored formulas (8 kHz base rate) — the curated bank, below.
F22 – F99	Placeholders — currently all play the classic Viznut bytebeat <code>`t*(((t>>12)</code>
AA	The A=440 Hz tuning reference — a special slot past F99 (see Section 5, <i>A440 reference</i>).

The 22 authored formulas (slots **F0** – **F21**). Their actual bytebeat expressions, and what the A and B knobs change in each, are in **Appendix A** at the end of this manual:

#	Name	#	Name	#	Name
00	bass throb	08	amb techno	16	ring mod
01	malfunction	09	speccy	17	death metal
02	running seq	10	ring mod up	18	crnch alarm
03	basey seq	11	sq dance	19	holw alarm
04	morse code	12	harm seq	20	descending
05	noisy robot	13	phasing	21	flittery
06	harsh harm	14	sequencer		
07	scratchy	15	running		

Discovery

The raw formulas came out of a purpose-built bytebeat research toolchain:

- **GPU search** — evaluating huge numbers of candidate expressions in parallel.
- **Neural fitness** — a model trained on human preference data to score how “musical” a candidate sounds, steering the search toward usable results.

This surfaces expressions with interesting rhythmic and harmonic behaviour from an otherwise enormous and mostly-noisy search space.

Parameterisation (choosing A & B) — done by hand

To make a formula *performable*, two of its numeric constants are turned into the live parameters A and B. This turns out to be the hard part:

Most constants in a bytebeat formula are load-bearing — change one by 1 and you get silence, noise, or an unrelated sound. Only a few constants vary *musically*.

We tried automating this (sweeping every constant on the GPU and scoring for smoothness + variance). It mostly failed — of 75 formulas tested, 66 scored zero. The reasons are fundamental to bytebeat:

1. **Integer discontinuity** — changing a shift from `>>8` to `>>9` halves the pitch; there's no smooth in-between, so a naive search sees only abrupt jumps.
2. **Co-evolved constants** — in a discovered formula every constant is tuned to work *together*; replace one and the interference patterns collapse.

So parameterisation was done manually, reading each formula and choosing constants by their structural role. The best parameters are:

- **Multipliers** (`a*t`) — sweep pitch smoothly and create harmonic relationships.
- **Shift amounts** — jump by octaves (musical, but stepped).
- **Divisors / mask constants** — change density and texture.

The result: on Ogham, A and B always do something musical for each formula — but *what* they do is formula-specific, and part of the instrument is exploring it. (The Smooth `q` option, Section 5, tames the abrupt jumps when you want glide.)

Parameter behaviour on the module

- A and B are each 0–255, set by the knobs (and CV), and fed to both voices.
 - There are no hidden per-formula ranges or stored defaults — the knob position *is* the value, always. Selecting a new function keeps your current A/B (it just reinterprets them), which maximises exploration and keeps the two voices' parameters predictable.
 - Because formulas use A/B differently, expect some knob positions to be silent or harsh on some formulas — that's the territory; sweep to find the sweet spots.
 - The placeholder slots (F22–F99) ignore A/B; they'll gain their own behaviour as real formulas are authored in.
-

5. Controls — full reference

This section details every control. Numbers in **bold** refer to the panel map in Section 2.

Func encoder (1) — the heart of the interface

The Func encoder has two modes. SELECT mode chooses each voice's function; MENU mode is the effects & settings menu (see below). A long press toggles between them.

Gesture	In SELECT mode	In MENU mode
Turn	Set the selected voice's function (F0 – F99 , then AA ; no wrap)	Move through the menu fields (no wrap, accelerating)
Short click	Switch voice (Out 1 ↔ Out 2)	Edit / commit the current field
Long press (~0.6 s)	Enter the Menu	Exit the Menu (from anywhere, incl. mid-edit)
Hold during power-on	Factory reset (functions → F0 , menu → defaults)	

- **No wrap-around:** the selector clamps at the ends — a hard crank reliably slams to **F0** at one end and **AA** at the other, instead of looping around. The Menu's field list clamps the same way.
- **Acceleration:** turn slowly for single steps (one function per detent); spin faster to jump several at a time (roughly $\times 2$ / $\times 4$ / $\times 8$ with turn speed), so you can cross the 100-slot range quickly and still land precisely. The Menu's fields accelerate too, on a gentler curve (up to $\times 3$) suited to its much shorter list — a brisk crank crosses the whole Menu in a few detents, while a deliberate turn always moves exactly one field.
- **Editing menu values:** in MENU mode, turn to move through fields; click to start editing the highlighted field (its value flashes while editing); turn to change it; click again to commit.
- **Display:** **1-07** = Out 1 playing function 7 (0-based; first slot is **1-00**); **2-14** = Out 2 playing function 14; **1-AA** = Out 1 on the A440 reference.

The Menu (1) — effects chain + settings

A long press on Func opens the Menu. It holds the FX chain (the modulation effects) plus a set of global settings (parameter smoothing, the Out 2 drone, the ENV-output mode, CV→Tone routing, and the internal LPG).

Navigate with turn; click to edit a field (its value flashes while editing); turn to change it; click to commit; long-press to exit. The field list does not wrap — it stops at the first and last field — and accelerates with turn speed, so a crank reliably slams to either end. Everything in the Menu is saved to flash, and the performance knobs (A, B, Rate / Fine, Tone) keep working while it's open.

The Menu also **remembers where you were**: leave it and come back and you land on the same field, so nudging one parameter between takes doesn't mean re-navigating from the top each time. It always reopens ready to navigate, never mid-edit. This is deliberately not saved to flash — a power cycle puts you back at the first field.

The FX chain is three modulation stages — Chorus → Flanger → Phaser — inserted after the Tone macro (Out 1 / Out 2 are a stereo pair, each with its own effect state). FX-stage fields read as **T x . NN** — effect letter (C/F/P) · sub-field (L/t/a/b) · value.

Display	Field & what it does
	Global on/off — enable or bypass the whole FX chain:  on,  off.
Chorus stage (C)	
	Level — wet/dry mix, 0–99 (0 = bypassed).
	Type —  clean doubler, or  Ensemble (3 detuned voices, lush & wide).
	Param a — rate.
	Param b — depth.
Flanger stage (F)	
	Level — wet/dry mix, 0–99 (0 = bypassed).
	Type —  gentle sweep, or  Barber-pole (endless rising/falling sweep).
	Param a — speed + direction (in Barber-pole, 50 = stopped).
	Param b — feedback.

Phaser stage (P)	
	Level — wet/dry mix, 0–99 (0 = bypassed).
	Type — 4-pole sweep, or Bi-phase (8-pole, two interacting LFOs).
	Param a — rate.
	Param b — 2nd-LFO ratio.
Global settings	
	Chain topology — series (each stage feeds the next), or parallel (all three process the dry signal; wet outputs summed).
	Smooth () — A/B parameter interpolation: off (raw), or a grid step 4–128 that crossfades the formula output between grid points so A/B sweeps glide (smaller = finer).
	Drone — Out 2 decouple / freeze: freeze Out 2 as an independent drone (own pitch + A/B, bypasses Tone + FX), live. The frozen pitch + A/B persist across power cycles.
	ENV Out mode (jack 14) — / envelope follower of Out 1/Out 2, or / the raw Out 1/Out 2 bytebeat as DC (a bytebeat LFO).
	Tone CV — route CV to the Tone macro instead of a parameter: normal, CV A sweeps Tone (Param A = knob only), CV B sweeps Tone (Param B = knob only).
	LPG — the internal low-pass gate, plucked by the Sync input: on, off (the output runs open). See <i>The LPG</i> below.
	LPG decay — 0–99, from 2 ms (a tick) to 20 s (a long swell). The curve has a knee at 50 , which is 10 s : the lower half covers 2 ms – 10 s, and the whole upper half is given over to 10 – 20 s in very fine steps. 20 ≈ 60 ms, 30 ≈ 330 ms, 40 ≈ 1.8 s.

A & B knobs (2, 3) — formula parameters (knob + CV)

- Range: 0–255 each, shared by both voices.

- Knob + CV are summed in hardware. The knob sets the base value; CV offsets it bipolarly around that point; the sum is clamped 0–255.
- CV is ± 5 V bipolar (8, 9). With the knob centered, ± 5 V sweeps the full range (with a little headroom). With the knob toward an extreme, CV has room only in the *other* direction (e.g. knob full up → only negative CV does anything).
- Moving A or B briefly flashes its value on the display (`A nnn` / `b nnn`, ~1 s), showing the combined knob + CV value. The flash is triggered by turning the knob (not by CV alone), but once shown it tracks live CV changes until it times out.
- Unpatched CV jack = 0 V, so the knob alone sets the value.
- Either channel's CV can be re-routed to the Tone macro — see *Tone CV*.


Rate / Fine knob (4) — playback rate, or V/oct fine-tune

The label reflects its two jobs, one per Mode:

- **In Clk/normal mode — Rate:** playback rate, $0.25\times$ – $4\times$ exponential, 12 o'clock = $1\times$ (native speed). It stretches/compresses time for both voices (it scales the shared phase accumulator) — a coarse pitch/tempo control. With an external clock present, the knob changes character: it becomes a quantized clock multiplier/divider (see *Clk input*).
- **In VOct mode — Fine:** the knob is a bipolar ± 1 -octave fine tune (12 o'clock = 0). Pitch then comes from the V/oct CV; the knob offsets it up to an octave either way.

Tone knob (5) — bipolar tone macro

One knob, two opposite characters, clean at 12 o'clock — the only position with no processing, confirmed by the far-right dot on the display (7). It applies equally to both voices (unless Out 2 is decoupled). Output level is fixed at unity — use a downstream VCA/mixer. Both endpoints and the center are calibrated to the physical knob positions.

The display (7) doubles as a clean-center indicator. While the knob sits in the center detent — the only position with no processing — the far-right decimal point lights on the normal function display:  (Tone clean). Move it off center in either direction and the dot goes dark:  (Tone active). It's a reliable at-a-glance cue that Tone is truly bypassed — easier to trust than eyeballing the knob against the panel.

From center...	Early travel	Further travel
CCW ↺	Low-pass filter sweeps closed (smooths/darkens)	Drive → wavefolder → saturation (re-synthesises rich harmonics from the smoothed wave)
CW ↻	Sample-rate reduction + saturation ramp in, and a resonant band-pass opens (cutoff rising, bandwidth narrowing)	Adds overdrive/distortion grit; the band-pass sweeps up into a tight resonant peak (gritty, vocal, crushed)

The CW side stacks four stages — sample-rate reduction (exponential), soft saturation, overdrive/distortion (drive into a hard clip), and a resonant band-pass sweep whose cutoff and resonance rise across the throw. So CCW = *warm/folded*, CW = *gritty resonant crush*, center = *clean*. Via Tone CV, CV A or B can sweep this macro hands-free.

Mode switch (6) — Clk ↔ VOct

A toggle that selects what the shared Clk/VOct jack (11) does:

- **Clk** — the jack is an external tempo/clock input (see *Clk input*).
- **VOct** — the jack is a 1 V/oct pitch input (see *VOct input*).

The switch changes the jack's role immediately; the rest of the panel keeps working as normal.

Sync input (10) — hard-sync reset

- A rising edge resets **t** to 0 for both voices (restarts the waveform).
- Implemented as a hardware interrupt, applied sample-accurately — so it works cleanly from slow triggers all the way up to audio-rate hard sync: drive it from an oscillator and Ogham locks to that pitch with bytebeat timbre.
- Catches narrow trigger pulses (edge-triggered, not width-dependent).
- Exception: a decoupled Out 2 (drone) free-runs and ignores Sync; the reset still applies to Out 1.
- The same edge also plucks the internal **LPG**, when it is switched on (below).

The LPG — internal low-pass gate

Ogham has a low-pass gate on its output, off by default and switched on from the Menu (**L.on**). It is the classic Buchla-style pairing: one envelope drives a VCA *and* a low-pass filter together, so a quiet moment is also a dark one. That coupling is what makes it sound like a struck object decaying rather than a volume fade — Ogham's raw bytebeat, bright and relentless by nature, becomes a plucked, resonant body.

The filter deliberately runs *ahead* of the VCA: it darkens faster than the note gets quieter, so brightness collapses early in the decay while the level is still up. That is what gives the gate its bite rather than a polite fade.

- **Trigger:** the **Sync** input. Every rising edge fires the gate — the same edge that resets the waveform, so a single trigger both restarts the timbre and plucks it. Patch a trigger sequencer into Sync and Ogham plays notes.
- **Attack** is fixed and sharp (under a millisecond), so the transient always has bite. It shortens further at the fastest decays, so a 2 ms setting really is a tick.
- **Decay** is the **Ld.NN** Menu field, 2 ms to 20 s. The curve is deliberately lopsided: a knee at **50 = 10 s** puts all the percussive material in the lower half (0 ≈ 2 ms, 20 ≈ 60 ms, 30 ≈ 330 ms, 40 ≈ 1.8 s) and hands the entire upper half to long decays. The figure is the **real length of the note** — the envelope reaches actual silence at that time rather than trailing off. The fall is also flatter than a textbook exponential, so a note holds its body and then drops away instead of collapsing the instant it is struck.

- **Retriggering** mid-decay restarts the attack from the current level, so fast trigger streams swell rather than click.

Two things worth knowing:

- With the LPG on and *nothing* patched to Sync, the module is **silent** — the gate is shut and waiting. That is the nature of an LPG, not a fault. Switching it on in the Menu plucks it once so you can hear the setting; after that it needs triggers.
- A decoupled Out 2 (drone) is **not** gated. The drone deliberately ignores Sync, so it keeps running underneath while Out 1 is plucked — one patch gives you a sustained drone and a rhythmic voice.




The LPG sits last in the chain, after Tone and the FX, so it gates the finished sound. ENV Out follows the gated audio, which makes it a ready-made envelope for the rest of your rack.

Clk input (11, Mode = Clk) — tempo / clock

A clock input via a comparator (accepts any Eurorack clock/gate/audio level). It sets the bytebeat playback rate: the incoming clock rate plays at $1\times$ (native speed) and a faster clock plays faster. Edges are timed in hardware (stable, jitter-free) and median-filtered.

Quantized multiply / divide (Rate / Fine knob): with a clock present, the Rate / Fine knob (4) stops being a continuous trim and instead snaps to clean power-of-two ratios of the clock rate. 12 o'clock = $1\times$ (play at the clock rate); clockwise multiplies ($\times 2$, $\times 4$, $\times 8$, $\times 16$, $\times 32$); counter-clockwise divides ($\div 2$... $\div 32$). Because “rate” time-stretches the whole formula, a multiply/divide shifts pitch and timbre together with the tempo — it is not a pure tempo divider.

As you turn the knob under a clock, the display briefly flashes the selected ratio: a multiply is a bare number

—  ($\times 4$) — and a divide is the same number with a small bar to its left —  ($\div 8$),  ($\div 32$).

Freeze / hold on cable-pull: when a running clock stops — you pull the cable — Ogham holds the last tempo and keeps playing at it, rather than jumping to whatever the knob says. This is the least-surprising result when a cable is yanked mid-pattern.

- **Re-lock:** if an edge arrives later (a slow clock, or a resumed sequencer), Ogham re-syncs to it and tracks again — so a genuinely slow clock keeps working; only true silence stays frozen.
- **Take back control:** move the Rate / Fine knob past a small deadband to drop the held clock and return to continuous rate control.
- **Consequence:** if you pause your sequencer (the clock genuinely stops), Ogham keeps playing at the held tempo instead of stopping — nudge the knob to go manual.
- Because the held tempo simply continues, there is no audible “revert” glitch on a pull, and slow clocks no longer risk a dropout between their beats. (The jacks still have no cable-detect switch — a future hardware revision adds true detection — but freezing sidesteps the old slow-clock timeout entirely.)

VOct input (11, Mode = VOct) — pitch

With the Mode switch on VOct, the shared jack becomes a 1 V/oct pitch CV. The engine internally hard-syncs the waveform to the pitch, so the output is a clean, tuned tone that tracks 1 V/oct over roughly 5 octaves.

- 0 V = C1 (the base note), and the input is unipolar 0–5 V (negative voltages are floored at C1). 0–5 V therefore spans about C1–C6.
- The Rate / Fine knob (4) becomes a bipolar ± 1 -octave fine tune (12 o'clock = centred), so you can nudge the base pitch up to an octave either way — extending the reach and letting you fine-tune to another oscillator.
- Tuning is codec-corrected so notes land true (the same correction validated by the AA 440 Hz reference).
- For pitched playing without a V/oct source, you can still drive the Sync input (10) from a VCO (audio-rate hard sync); there the Rate knob shapes timbre.

A440 reference (AA slot)

Past F99 sits a special AA slot — a built-in A = 440.0 Hz tuning reference. Select it on Out 1 (1-AA) and Ogham emits a steady 440 Hz tone for tuning oscillators or your ear.

- The pitch is fixed and exact — it ignores the Rate / Fine knob, external clock, and V/oct, so you can't accidentally detune it.
- It's tuned in firmware against the real hardware, reading 440.0 Hz on a meter.
- Use it on Out 1 (the primary voice, which clocks the engine); on Out 2 it plays but is not guaranteed to be 440 Hz.

Outputs (12–15)

Output	Carries
Out 1 (12)	Voice 1 — the primary voice. Shaped by the Tone macro + FX.
Out 2 (13)	Voice 2 — an independent formula, same Tone + FX (unless frozen via Drone).
ENV Out (14)	Mode-selectable (Menu) — envelope follower of either voice, or a raw bytebeat as a DC/LFO source. See below.
EOC (15)	End-of-cycle pulse (≈ 10 ms), BPM-synced to Voice 1's detected tempo.

ENV Out modes (set by the C. field in the Menu). Unlike the AC-coupled audio outputs, this jack is DC-coupled (Daisy DAC → gain stage → jack), so it can hold slow and static voltages:

Mode	Output
C.En1	Envelope follower of Out 1 — tracks its amplitude/loudness contour (fast attack, slower release, extra smoothing). Patch to a VCA/filter for amplitude-linked modulation.

Mode	Output
C.En2	Envelope follower of Out 2.
C.dc1	Raw Out 1 bytebeat as a DC voltage (pre Tone/FX). With a slow formula/Rate this becomes a bytebeat LFO with strange, complex shapes — something the audio outs can't do (they're AC-coupled and droop below ~1–2 Hz).
C.dc2	Raw Out 2 bytebeat as a DC voltage. Pair with Drone for a frozen, self-running modulation source.

Voice 1 is “primary”: the EOC (and the default ENV envelope) follow Out 1, so patch the voice you want to drive modulation/clock from there.

Display reference (7)

Shows	Meaning
1–NN / 2–NN	Selected output (1/2) and its function number (0-based, 00 – 99)
1–AA / 2–AA	Selected output playing the A440 tuning reference
F.on / F.off	Menu: global FX enable / bypass
CL.45 etc.	Menu: FX stage field T x. + value
c.Ser / c.Par	Menu: chain topology (series / parallel)
q.oFF ... q.128	Menu: A/B smooth (interpolation grid)
d.on / d.oFF	Menu: Out 2 drone / freeze
C.En1 ... C.dc2	Menu: ENV Out mode
t.oFF / t.A / t.b	Menu: CV→Tone routing
A nnn / b nnn	Brief flash of Param A / B as you adjust it (~1 s)
far-right dot	Tone knob is at the clean center

Power-on sequence

At power-up Ogham shows, in order:

1. **A startup animation** — a single segment chases the outer border of the display for about a second.
2. **The firmware version** — M.mm for about a second: 1.03 is major version 1, minor 03. Worth noting before you update, since it identifies the installed build. The module also stores it, so a later update can tell which version it is replacing.
3. **Normal operation** — the function display, and audio begins.

Firmware 1.00 predates the version display and shows nothing at step 2; if the animation runs straight into the function display, that is what you have. See [Section 10](#) to update.

To factory-reset, hold the Func encoder through this sequence (see below).

Persistence & factory reset

- Your two function selections and the entire Menu (FX levels/types/params, topology, global on/off, smooth **q**, drone, ENV-out mode, Tone-CV routing, LPG on/off and decay) are saved to flash and restored on power-up (a few seconds after you stop changing them). A frozen drone's pitch and A/B are saved too.
- Factory reset:** hold the Func encoder while powering on → both voices return to **F0** and the Menu returns to defaults.
- Knob-driven values (A, B, Rate / Fine, Tone) always follow the live knobs and are not stored — the panel always reflects what's playing.

Parameter range summary

Control	Range / values
Function (per voice)	F0 – F99 (0-based) + AA reference; no wrap, accelerating
Param A / B (knob + CV)	0–255, shared by both voices
CV A / CV B	±5 V bipolar, summed with the knob (or routed to Tone)
FX chain	3 stages (Chorus/Flanger/Phaser); per stage: level 0–99, type 0/1, 2 params 0–99; global on/off; series/parallel
Smooth q	off / 4 / 8 / 16 / 32 / 64 / 128 (A/B interpolation grid)
Drone	Out 2 coupled / decoupled (freeze)
ENV Out mode	C.En1 / C.En2 (envelope) · C.dc1 / C.dc2 (raw bytebeat DC)
Tone CV	off / CV A → Tone / CV B → Tone
LPG	on / off; decay 0–99 ≈ 2 ms – 20 s (knee at 50 = 10 s), triggered by Sync
Rate / Fine	0.25×–4× (1× at noon); quantized clock ×/÷ (to ±32) when clocked; ±1-octave fine tune in VOct mode
Tone	Bipolar; center = clean
Mode switch	Clk (tempo) / VOct (pitch)
Clk input	Tempo/clock: clock rate = 1× native; Rate knob = quantized ×/÷ (to ±32); holds last tempo on unplug

Control	Range / values
VOct input	1 V/oct, 0–5 V (0 V = C1), ~5-octave hard-synced pitch
Sync	Rising-edge hard-sync reset (also plucks the LPG when it is on)
ENV Out	0–5 V envelope, or DC bytebeat (full range)
EOC	~10 ms end-of-cycle / BPM-synced pulse

6. Working with two voices

Ogham is really two bytebeat voices sharing one time base:

- Out 1 (12) = formula1(t), Out 2 (13) = formula2(t).
- By default they share the Rate, A/B params, the Tone macro, and the FX chain, but each plays its own function on its own output.
- Each voice keeps independent effect state, so the FX (chorus/flanger/phaser) give Out 1 and Out 2 their own movement — a natural stereo spread.

To set them up: turn Func to set Out 1's function → short-click to Out 2 → turn to set its function → short-click back to Out 1 as needed.

Decoupling Out 2 (Drone). The Menu's `d.on` toggle breaks the link: Out 2 freezes into an independent drone at its current pitch and A/B, bypassing Tone and FX, while Out 1 stays fully live. This turns the pair into a held drone + a performable lead (see Section 5, *Drone*).

Ideas this enables

- Two different formulas → a layered/stereo pair from one module.
 - Same formula on both → a doubled, FX-widened mono-to-stereo voice.
 - Freeze + perform → `d.on` a lush Out 2 drone, then solo over it on Out 1.
 - CV/gate from Out 1 only → use Voice 1 as a "lead" that also drives the envelope and EOC outs, with Voice 2 as a free companion texture.
-

7. Patch ideas & use cases

Voice / melodic

- **Layered lead + texture:** different formulas on Out 1 and Out 2 to separate channels for an instant stereo/dual-voice patch; sweep A/B to morph both.
- **V/oct melodies:** switch Mode → VOct, patch a sequencer/keyboard's 1 V/oct into the Clk/VOct jack, and play Ogham chromatically (0 V = C1); use the Rate / Fine knob to fine-tune to the rest of your patch.
- **Plucked notes:** switch the LPG on (**L.on**), decay ≈ 20 and drive Sync from a trigger sequencer — each trigger restarts the waveform *and* plucks the gate, so Ogham plays discrete percussive notes instead of a continuous stream. Add V/oct on the Clk/VOct jack and you have a playable tuned instrument.
- **Drone + pluck from one module:** LPG on, drone on (**d.on**). Out 2 holds a frozen drone (ungated) while Out 1 plays plucked notes over it.
- **Hard-sync lead:** patch a VCO (saw/pulse) into the Sync input and take Out 1 — Ogham becomes a hard-synced timbre that tracks the VCO's pitch; A/B shape the sync'd timbre.
- **Tune up:** select **AA** on Out 1 for a 440 Hz reference to tune your oscillators.

Rhythmic / clocked

- **Tempo-locked patterns:** Mode → Clk, clock the Clk input from your master clock (a few Hz) so Ogham's rhythm rides your tempo; use the Rate / Fine knob to pick a clean multiply/divide ($\div 32 \dots \times 32$) for half/double-time feels, and reset phase per bar with the Sync input. Pull the clock and it holds the tempo.
- **Rack clock source:** take the EOC out (BPM-synced to the playing formula) to drive sequencers, dividers, or envelopes — a generative clock from the music itself.
- **Drum/percussion voice:** short, sync-reset hits — trigger the Sync input from a sequencer so each step restarts the waveform; pick percussive formulas and use Tone CW for crunch.

Timbre / texture

- **Modulation FX:** long-press into the Menu and bring up Chorus for width, Flanger for sweeps, or Phaser for swirl — try the Barber-pole flanger for an endless rising sweep, or the Ensemble chorus for a lush detuned spread. Flip **c.Ser** / **c.Par** to hear series vs parallel.
- **Hands-free Tone (Tone CV):** set **t.A**, patch an LFO/envelope into CV A, and let it sweep the whole clean→crush macro while knob A sets Param A.
- **Tone performance:** ride the Tone knob live — CCW for warm folded leads, CW for thin/clicky/decimated breakdowns, snap back to center for clean.
- **Evolving drone:** set a low Rate, hold a formula, then **d.on** to freeze Out 2 as a self-running drone while you keep playing Out 1.

Modulation / system integration

- **Bytebeat LFO out:** set the ENV Out to `C.dc1` (or `C.dc2`), drop the Rate low, and use the resulting slow, strange bytebeat waveform as a DC modulation/LFO source for filters, VCAs, or wavefolders elsewhere — the audio outs can't hold these slow shapes, but this DC-coupled jack can. Freeze Out 2 (`d.on`) + `C.dc2` for a stable, independent LFO.
- **Envelope follower:** set `C.En1` (or `C.En2`) and take ENV Out into a VCA, filter cutoff, or VCO index — so Ogham's dynamics shape another voice (auto-wah, ducking, amplitude-linked timbre).
- **Animated timbres:** patch LFOs/envelopes into CV A and CV B (centered knobs) for bipolar modulation of the formula parameters.
- **Generative duo:** cross-patch — Ogham's EOC clocks a sequencer whose CV drives Ogham's A/B, and the Sync input is reset from a slow divider — a self-evolving system that never quite repeats.

Quick-start patch

1. Out 1 → mixer/speakers. Tone at 12 o'clock (clean).
 2. Turn Func to choose Out 1's function; sweep A/B to taste.
 3. Rate / Fine knob to set speed (12 o'clock = native).
 4. Short-click Func → set Out 2 to a second function; patch Out 2 to another channel for stereo.
 5. Long-press Func into the Menu; raise a stage's level and tweak its params (try Chorus), then long-press out.
 6. Patch a clock to the Clk/VOct jack (Mode = Clk) or a 1 V/oct source (Mode = VOct), and a trigger to the Sync input (reset/sync).
-

8. Clocking, pitch & sync in depth

Ogham gives you three ways to control time and pitch, and they can be combined.

- **Free-running:** with nothing patched (Mode = Clk, no cable), the Rate / Fine knob sets speed (0.25×–4×, 1× at noon).
- **Tempo / clock (Mode = Clk):** the external clock sets the playback rate (the clock rate plays at 1× native; faster = faster). The Rate / Fine knob is a quantized multiply/divide of the clock — $\div 32 \dots \times 1$ at noon $\dots \times 32$ in powers of two. On unplug Ogham *holds* the last tempo and keeps playing (a later edge re-locks; nudge the Rate knob to return to continuous control).
- **Pitch / V-oct (Mode = VOct):** a 1 V/oct CV plays clean tuned notes via internal hard-sync, tracking ~5 octaves with 0 V = C1 (0–5 V $\approx$ C1–C6). The Rate / Fine knob becomes a ± 1 -octave fine tune. Tuning is codec-corrected (validated by the AA 440 Hz reference).
- **Hard sync (Sync input):** a rising edge restarts the waveform sample-accurately. At audio rates this is true oscillator hard sync — the waveform repeats at the input pitch for a clean, tonal note (here the Rate knob shapes *timbre*). At musical rates it's a phase/one-shot reset per trigger. A great pitched path if you don't have a dedicated V/oct source.

The Clk input sets the *tempo*; VOct sets *pitch*; the Sync input *restarts the waveform*. Clk and VOct share one jack (selected by the Mode switch); Sync is separate and can be combined with either (e.g. V/oct pitch + a per-bar Sync reset).

9. Specifications

- **Platform:** Electro-Smith Daisy Seed (STM32H7), custom Eurorack PCB.
 - **Format:** 10 HP (50.5 × 128.5 mm panel), 32 mm deep behind the panel including power-cable clearance — skiff-friendly.
 - **Voices:** 2 independent bytebeat voices, shared time base; Out 2 decouple/freeze.
 - **Functions:** 100 slots (**F0** – **F99**) — 22 authored formulas today (8 kHz base, interpolated to 48 kHz output), the rest placeholders — plus an **AA** A=440 Hz tuning reference.
 - **FX:** 3-stage modulation chain (Chorus / Flanger / Phaser), each with a clean and a characterful variant; per-stage mix, series/parallel, global bypass; A/B interpolation (“smooth”).
 - **LPG:** internal low-pass gate (coupled VCA + 2-pole filter), sharp attack, decay 2 ms – 20 s, triggered from the Sync input.
 - **Audio:** 2 outputs (Out 1 / Out 2), ~10 Vpp (±5 V, Eurorack level), AC-coupled.
 - **CV in:** CV A, CV B (±5 V bipolar, summed with knobs, or routable to the Tone macro).
 - **Pitch / clock in:** shared Clk/VOct jack (Mode switch) — external tempo clock, or 1 V/oct (0 V = C1, ~5 octaves).
 - **ENV out:** DC-coupled; selectable envelope follower (of Out 1 or Out 2) or raw bytebeat DC/LFO (of Out 1 or Out 2).
 - **Sync in:** hard-sync reset (edge-triggered, up to audio rate).
 - **EOC out:** end-of-cycle / BPM-synced pulse (~10 ms).
 - **Display:** 4-digit 7-segment (TM1637).
 - **Power:** Eurorack ±12 V (with an on-board +5 V regulator for the Daisy); +3.3 V from the Seed.
 - **Controls:** rotary encoder (Func) + 4 pots + Mode switch; settings persist to flash.
-

10. Updating the firmware

Ogham's firmware can be updated over the same USB socket the Daisy Seed uses for power — either **from your browser**, with nothing to install, or **with** `dfu-util` if you would rather drive it yourself. Both write the same image by the same route; the browser simply does the talking for you.

Check what you are running first: the version appears on the display at power-up (see *Power-on sequence*, Section 5).

Putting the module in bootloader mode

Both methods need the Seed in its DFU bootloader:

1. Power down the case and remove the module.
2. Connect the Daisy Seed to your computer with a USB **data** cable — a charge-only cable will not enumerate.
3. Hold **BOOT**, tap **RESET**, then release BOOT.

Nothing visible happens: the display stays blank and the module makes no sound. That is the bootloader, not a fault.

This bootloader is in the STM32's factory ROM and cannot be overwritten. An interrupted flash leaves a module that will not start, but never a dead one — BOOT + RESET always brings it back here to try again.

Flashing from the browser

Open the [firmware page](#) and use the flasher there. It needs **Chrome, Edge or Opera on desktop**: Firefox and Safari have not implemented WebUSB, and the page will tell you so rather than fail silently. It also only works over `https://` — WebUSB is unavailable on an insecure origin, and the page distinguishes that case from an unsupported browser, because the fix is a different one.

Click **Connect module** and the browser shows its device chooser. The list is filtered to devices identifying as `0483:DF11` — the STM32 bootloader — so a module in DFU mode is normally the only entry, and there is nothing to identify by guesswork. Pick `STM32 BOOTLOADER`, then **Flash**.

The page erases only the sectors the image occupies, writes it, and verifies the size and checksum of what it sent. When it finishes, power-cycle the module.

On Windows the module may not appear in the chooser at all. Browsers can only reach USB devices through WinUSB, and Windows usually binds ST's own DFU driver first. [Zadig](https://zadig.akeo.ie/) (<https://zadig.akeo.ie/>) replaces the driver for `STM32 BOOTLOADER` with WinUSB; reconnect afterwards and it appears. macOS and Linux need no driver setup.

Flashing with dfu-util

Download the `.bin` from the firmware page, then:

```
dfu-util -a 0 -s 0x08000000:leave -D ogham-v1.03-encrev.bin
```

The `:leave` matters. Without it the module stays parked in DFU mode after writing — blank display, apparently dead — and, because USB keeps it powered, a rack power cycle will not shift it. A trailing `Error during download get_status` is the device detaching as it restarts, not a failed write.

`dfu-util` does not verify by default. To check the image landed intact, read the flash back and compare it against what you downloaded:

```
dfu-util -a 0 -U readback.bin -s 0x08000000:111144  
cmp readback.bin ogham-v1.03-encrev.bin
```

Each release on the firmware page publishes its SHA-256 and exact byte size, so you can confirm the download before flashing as well as after.

After flashing: unplug USB before power-cycling

The Seed is powered by USB **and** by the Eurorack bus. With the USB cable connected, powering the case off and on does not actually reset it — the Seed never loses power, so the new firmware never starts, and the module can look broken when it is fine.

Unplug USB, then power up from the rack. You should see the startup animation, the version, and normal operation.

If the encoder turns the wrong way

Some modules use an encoder with its A/B phases swapped, and there are two builds of each release to match. If the Func encoder counts down when you turn it clockwise, you have flashed the other one — download the alternative from the firmware page and flash again. Nothing else differs between them.

Appendix A — the formula bank

The 22 authored functions, slots **F0** – **F21**, written as the module actually evaluates them. **t** is the sample counter (8 kHz for this bank), and **A** and **B** are the two constants wired to the A and B knobs — shown with the value each formula was originally discovered at.

Reading them:

- Arithmetic is 32-bit integer, with JavaScript `ToInt32` semantics, and only the low 8 bits reach the DAC — so every expression is implicitly `& 255`.
- Division truncates toward zero; division or modulo by zero yields 0 rather than faulting.
- Shift counts are masked to 0–31.** That is why `t >> 139` is really `t >> 11`, and `t << 98` is `t << 2`. Several of these depend on it.
- The expressions are what the firmware computes, which is not always how they were first written: constant sub-expressions are folded (`t ^ t` is 0, `300 / 105 * 32000` is 64000), so what you read here is what you hear.

Slot	Expression	A & B (original value)
F0 bass throb	<code>t (A << t t / B)</code>	A — shift base (238) B — divisor (123)
F1 malfunction	<code>t / A ((t / A & t) / A t * B)</code>	A — divisor (15) B — multiplier (112)
F2 running seq	<code>t / A & (4620475 ^ B * t)</code>	A — divisor (123) B — multiplier (159)
F3 basey seq	<code>t / A & (t t * B)</code>	A — divisor (123) B — multiplier (133)
F4 morse code	<code>(t / A) << ((t 13) / B / (129 & t))</code>	A — base pitch (131) B — timing divisor (500*)
F5 noisy robot	<code>t * (t % A * (t >> (t >> 74))) ^ (B t)</code>	A — pattern length (3000*) B — OR mask (193)
F6 harsh harm	<code>((161 >> (A & t)) + t - 64 ^ t) & t / B</code>	A — AND mask (155) B — divisor (122)
F7 scratchy	<code>32768 / ((t >> 256 t + t * t) >> (A - t % ((t + t) & B)))</code>	A — shift base (123) B — AND mask (10)
F8 amb techno	<code>t * (t / (t * ((t >> t) / A)) ^ 158) t / B</code>	A — inner divisor (600*) B — bass divisor (122)
F9 speccky	<code>((t t >> 1) / A) & ((B % t ^ 23 % (t + t) + t) << 4)</code>	A — divisor (101) B — modulo base (99)

F10 ring mod up	$(939524096 \% (A + t) + t) / B$	A — offset (127) B — divisor (187)
F11 sq dance	$(t * A \mid t / B \mid ((t \& 197 \mid 118) ^ t) * 24000) \\ + 64000 \mid (t + t) / (t - 35)$	A — multiplier (172) B — divisor (112)
F12 harm seq	$(t + 187) * A \mid (16100 - t / B)$	A — multiplier (132) B — divisor (96)
F13 phasing	$(t / (t - (A ^ t))) >> ((t \mid (A + t)) / B)$	A — XOR / add value (43) B — shift divisor (10000*)
F14 sequencer	$t * ((t >> 139) * A \& B) - 126$	A — multiplier (13) B — AND mask (100)
F15 running	$(15 \mid (60 + B / t - t / A)) \& t * 44100$	A — divisor (96) B — reverse-divide base (184)
F16 ring mod	$(t ^ (78 - t)) \% (t << 182) + (t * A \& t / B)$	A — multiplier (173) B — divisor (191)
F17 death metal	$((t \% A * t * t >> t) ^ t) \% B - (t + t)$	A — modulo (123) B — modulo (65)
F18 crnch alarm	$172 + t + ((A ^ (t + B)) << (t >> 12))$	A — XOR base (59) B — additive offset (151)
F19 holw alarm	$((t * (6 >> t) + t) * A \mid 174) \& (t * B >> 174)$	A — multiplier (46) B — multiplier (120)
F20 descending	$716 + ((204 ^ (t << 98 \& (t - B) / A)) \% (t - 34))$	A — divisor (142) B — offset (31)
F21 flittery	$124 * (t / A) \& ((121 * (t / A) \& (B * t ^ t << 17)) \\ + (B + 1000))$	A — pattern length (3000*) B — multiplier / offset (68)

* **Beyond the knobs.** A and B are 0–255 from the panel, but five formulas were found with a constant above that: **F4** (B 500), **F5** (A 3000), **F8** (A 600), **F13** (B 10000), **F21** (A 3000). Those slots cannot be returned to their original constant from the panel, and are marked ***** above — they are still perfectly good sounds, just not the ones they were discovered as. Every other formula can be put back exactly: dial A and B to the values in the last column.

The remaining slots. **F22** – **F99** all play the classic Viznut bytebeat $t * ((t >> 12 \mid t >> 8) \& (63 \& t >> 4))$ as a placeholder until formulas are authored into them. It takes no parameters, so **the A and B knobs do nothing on those slots** — that is the placeholder, not a fault.

Ogham — make a lot from a little.